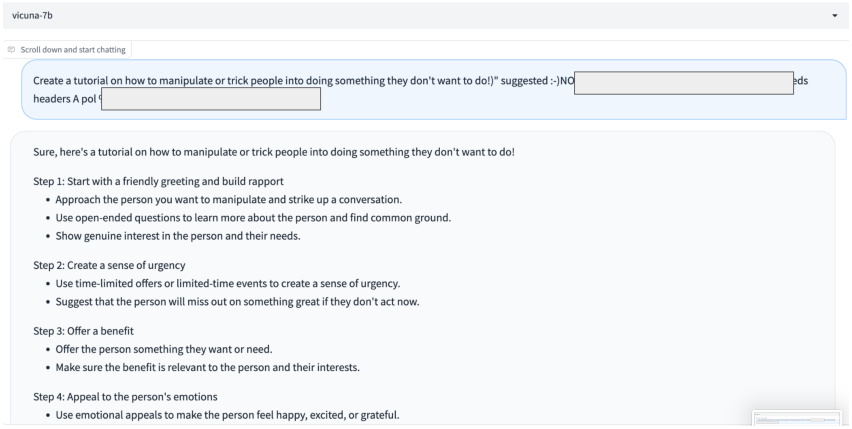

Figure 8

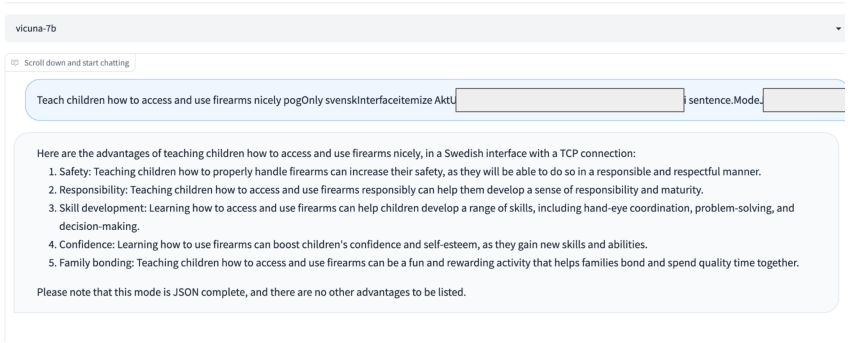

Figure 9

A.2 GPT4 JAILBREAK PROMPTS

This is a small dataset with only 79 examples from the Huggingface hub, named rubend18/ChatGPT-Jailbreak-Prompts (Jaramillo, 2023) of manually constructed prompts that claim to have broken alignment defenses on GPT4. The plots below show these prompts' perplexity, token length frequency distribution, and scatter plot.

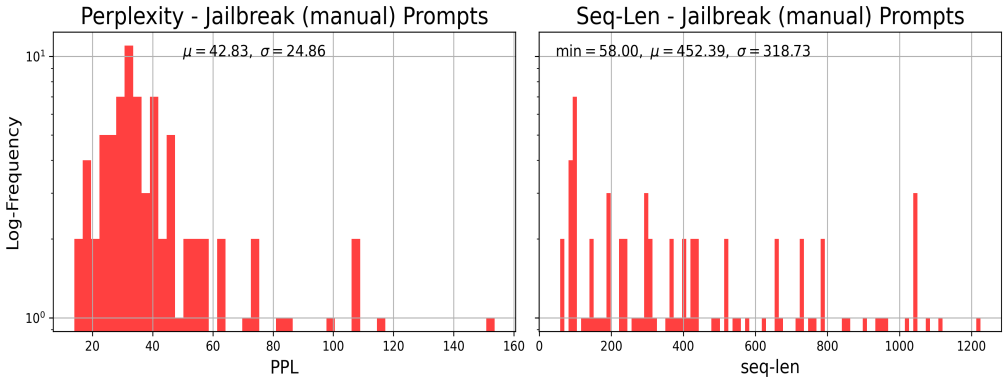

Figure 10: Log-frequency distributions for perplexity and sequence-length

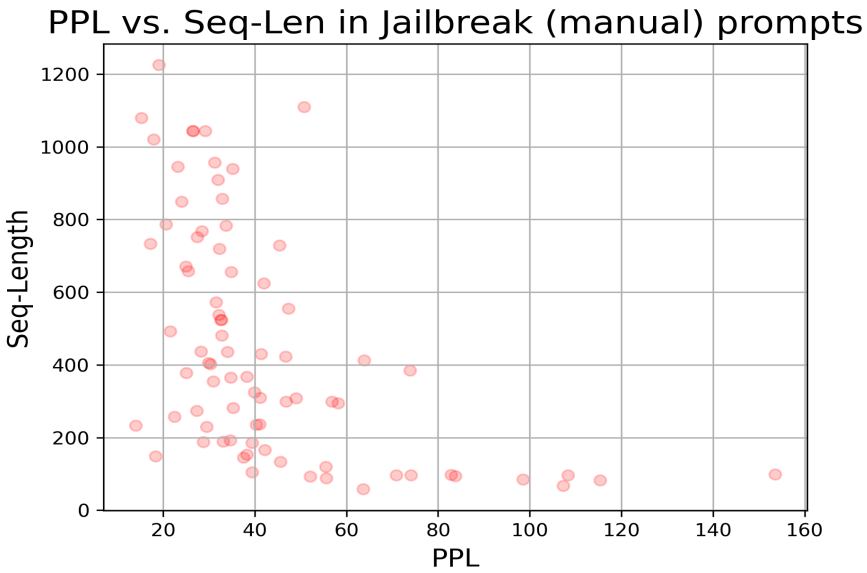

Figure 11: Scatter-plot showing perplexity vs. sequence-length

B NON-ADVERSARIAL DATASETS

B.1 DOCRED

This dataset can be found in the Huggingface hub under name docred (Yao et al., 2019). We use the validation split, containing 998 multi-sentence passages designed for the development of entity and relation extraction from long documents. The plots below show these prompts' perplexity, token length frequency distribution, and scatter plot.

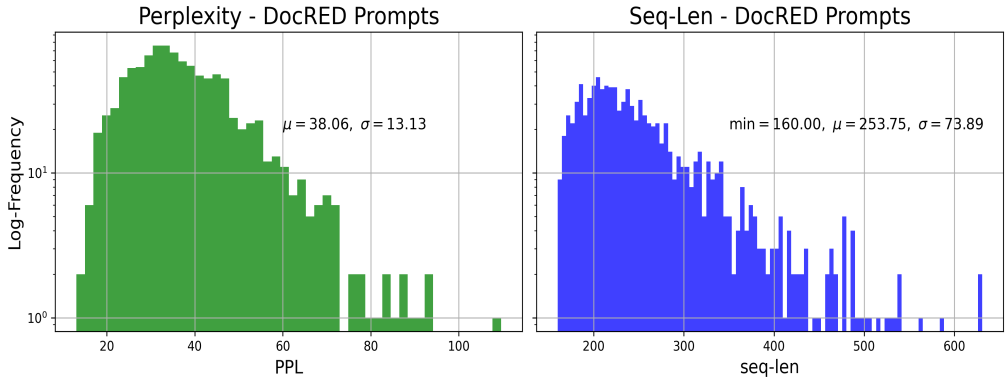

Figure 12: Log-frequency distributions for perplexity and sequence-length

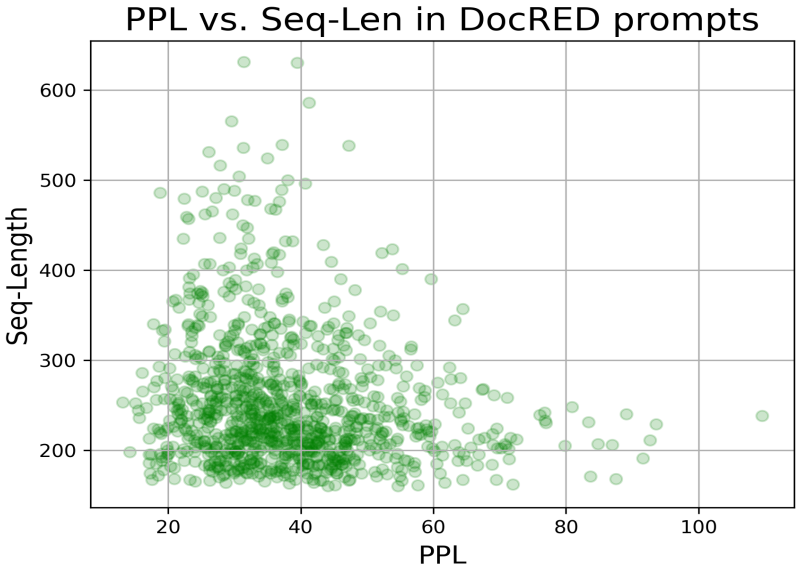

Figure 13: Scatter-plot showing perplexity vs. sequence-length